



VS.



APACHE HIVE

HIVE IS POWERFUL

- Makes MapReduce easy
- Easy to implement joins
- Optimize execution - partitions, buckets
- Support multiple file formats
- Schema on Read

```
select symbol, avg(volume) from stocks  
group by symbol
```



Hive vs RDBMS

- No pointed updates or deletes
- Limited indexing support
- High level transactional support*
- No triggers

APACHE HIVE

select symbol, avg(volume) from stocks
group by symbol



CLUSTER



VS.



Hive vs RDBMS

- No pointed updates or deletes
- Limited indexing support
- High level transactional support*
- No triggers

HIVE IS POWERFUL

- Makes MapReduce easy
- Easy to implement joins
- Optimize execution - partitions, buckets
- Support multiple file formats
- Schema on Read